
BigQuery vs. Cassandra

DBMS Implementation - Term Project
Winter '21 - CS587

Aaron Hudson
Travis McGowan

Benchmark Approach and Goals

- System Choices:
 - Wanted to explore systems new to us
 - Cassandra and BigQuery are very different from one another
 - Both effective for handling very large amounts of data
- Benchmark Approach:
 - Focus on clarity and simplicity with experimental design due to how fundamentally different the systems are
 - Include testing very large amounts of data
- Goals:
 - Determine cases where each DBMS shines
 - Explore areas each system is different and see how that impacts performance

BigQuery

- No Indices
- Made for BIG data
- Simple Materialized Views
- Multiple ways to load data
- Save and schedule queries
- Almost no customization options

Cassandra



Cassandra by Evelyn de Morgan
(1898, London); Cassandra in front
of the burning city of Troy
(<https://en.wikipedia.org/wiki/Cassandra>)

- No Joins
- NoSQL System (CQL)
- Indices: primary and secondary
- Employs a cluster of nodes to store data
- Performance is heavily driven by configuration

Testing Framework

4 Test Sets Comparing 3 Systems:

1. Selectivity
 - a. Percentage & Non/View
2. Scaling
 - a. Size & Selectivity
3. Tuple Insertion
 - a. Size & Non/View
4. Views

All Testing Performed on Google Cloud Platform
Virtual Machines

				BigQuery		Cassandra (No Index)		Cassandra (Indexed)	
				No views	<u>Views</u>	No Views	<u>Views</u>	No Views	<u>Views</u>
Selection	Percentage	(1, 10, 25, 50, 75)	selectivity	1	<u>2</u>	3	<u>4</u>	5	<u>6</u>
Use	Relations	MMTup, VMMTup_Current	selectivity						
Use	Relations	(OneK, TenK, CK, MM,XMM)	scaling	7	<u>8</u>	9	<u>10</u>	11	<u>12</u>
Selection	Percentage	(10, 50)	scaling						
Use	Relations	MMTup, VMMTup	updates						
Inserts	Distribution	Batch	updates	13	<u>14</u>	15	<u>16</u>	17	<u>18</u>
Selection	Percentage	(1, 5, 10, 25)	updates						

1. Scaling Experiment

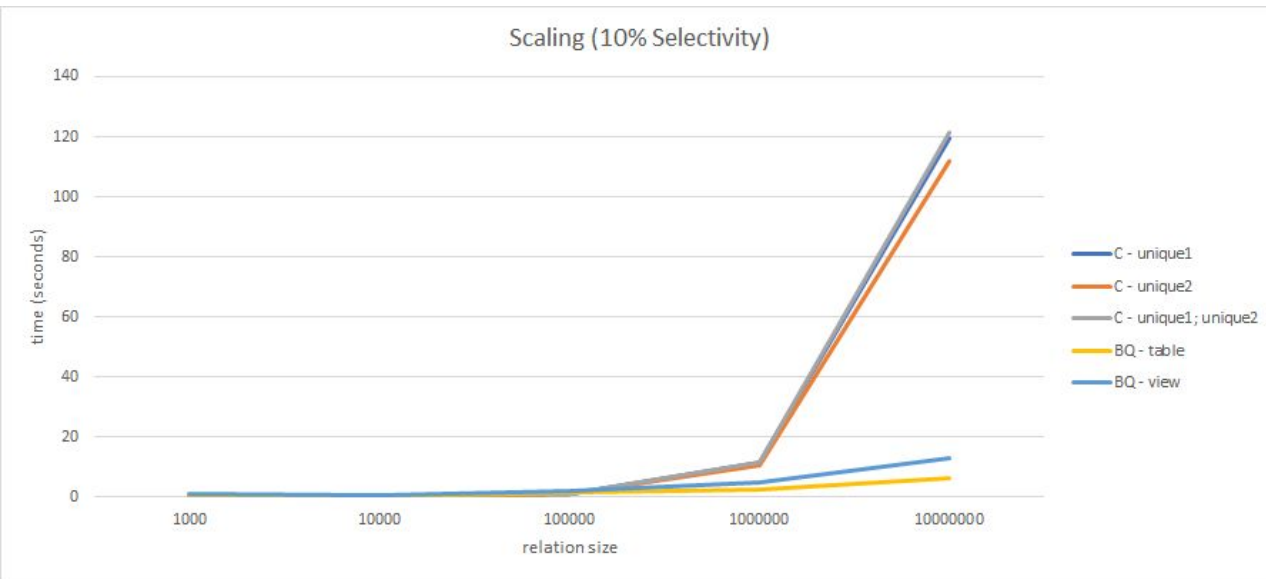
- Designed to compare how the systems handle varying data sizes (1k, 10k, 100k, 1mm, 10mm)
- Used queries with 10% and 50% selectivity
- General Query Format:

```
SELECT COUNT(unique2)
FROM relation_to_test
WHERE unique2 BETWEEN selectivity_range;
```

(Cassandra CQL doesn't have BETWEEN - used less than and greater than to replicate)

- Expected results:
 - As the relation size increases, BigQuery and Cassandra (no index) will degrade faster
 - Cassandra with an index will remain stable


1. Scaling (10%)



Size	1 k	10 k	100 k	1 mm	10 mm
C - U1	0.03	0.15	1.20	11.68	119.34
C - U2	0.02	0.12	1.08	10.77	111.71
C - U1/2	0.04	0.14	1.17	11.73	121.37
BQ - T	0.24	0.35	1.45	2.63	6.56
BQ - V	1.17	0.90	1.97	4.78	13.30

- BigQuery performs better on large relations
- BigQuery performs quite a bit better than this if you allow the results to come from the cache
- Relations of 1mm seems to be the beginning pivot point for Cassandra VM

1. Scaling (50%)



Size	1 k	10 k	100 k	1 mm	10 mm
C - U1	0.05	0.32	2.18	18.76	192.65
C - U2	0.06	0.25	2.19	17.67	182.71
C - U1/2	0.04	0.24	2.23	18.56	195.90
BQ - T	0.27	1.47	1.82	4.29	9.47
BQ - V	0.99	1.61	2.41	6.25	15.31

- Cassandra performance stays relatively stable until 100k records (as compared to 10% scaling)
- BigQuery performance noticeably slows around 10k

1. Scaling Experiment - Conclusions

- BigQuery:
 - Performance remained stable as relations increased in size
- Cassandra:
 - Performance remained stable from 1k - 100k relation sizes
 - At 1 million and beyond performance dropped significantly
 - Likely due to data being spread across the cluster of nodes ineffectively
- Actual Results (in contrast with Expected):
 - As the relation size increases, BigQuery remained stable
 - Cassandra with an index performed better than without one, but still had considerable performance degradation after 1 million tuples

2. Selectivity Experiment

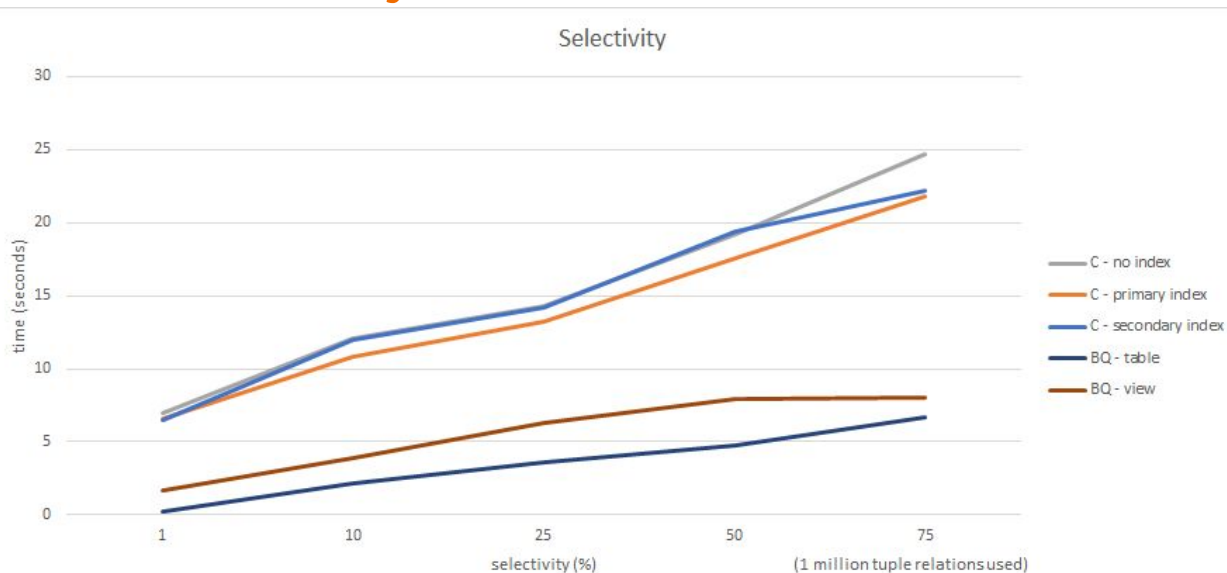
- Designed to compare how the systems handle queries of increasing selectivities (1%, 10%, 25%, 50%, 75%)
- Used relation with one million tuples for all queries
- General Query Format:

```
SELECT COUNT(unique2)
FROM MMTup
WHERE unique2 BETWEEN selectivity_range;
```

(Cassandra CQL doesn't have BETWEEN - used less than and greater than to replicate)

- Expected results:
 - At higher selectivities BigQuery and Cassandra without an index will perform better (10% rule)
 - At lower selectivities, Cassandra with an index will perform better

2. Selectivity



Selectivity	1%	10%	25%	50%	75%
C - no index	7	12.10	14.3	19.25	24.74
C - primary index	6.5	10.80	13.20	17.53	21.81
C - secondary index	6.5	12	14.23	19.41	22.24
BQ - table	0.2	2.11	3.61	4.72	6.67
BQ - view	1.6	3.92	6.26	7.93	8.07

Cassandra Performance:

- Overall much worse than BQ
- Primary index performed best
- Secondary index made little difference except at 75% selectivity

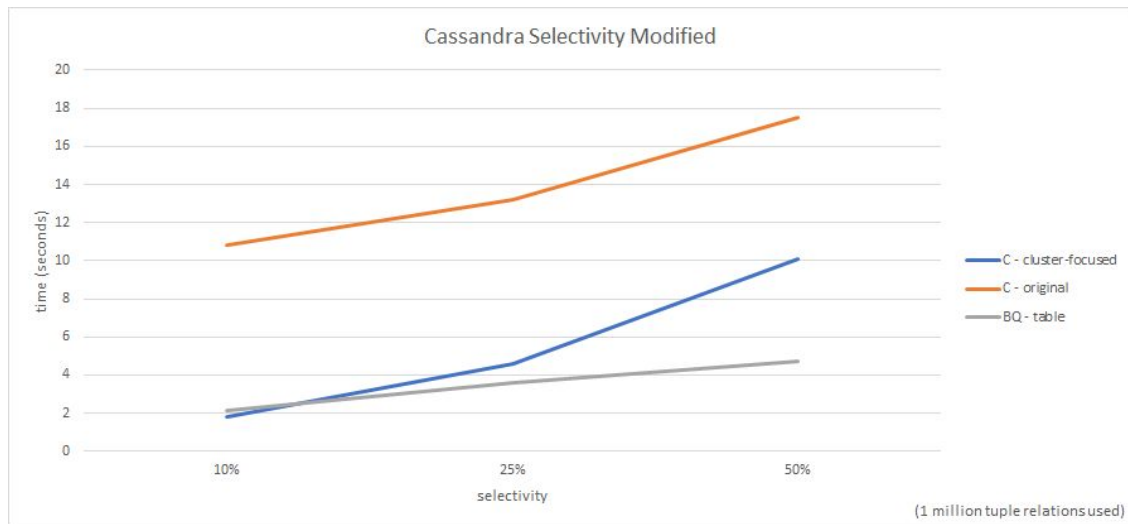
BigQuery Relations:

- Performed significantly better
- Use of materialized view hindered performance

Improving Cassandra Performance

- First attribute of Primary Key acts as a Partition Key
(determines how tuples are organized in nodes across the cluster)
- Goal is to query data in a way that stays within a node
- Original queries not great for this
 - Required ALLOW FILTERING setting - visited multiple nodes
- Need to build database conscientiously
 - How data will be stored across the cluster
 - How data will be queried

2. Selectivity - Cassandra Modified



- Original relations had Partition Keys of unique1 or unique2
- Cluster-focused relations have Partition Keys on ten, four, and two depending on desired selectivity
- Performance closer to BQ
(50% selectivity possibly requires visiting multiple nodes, which indicates the non-linear increase)

2. Selectivity Experiment - Conclusions

- BigQuery:
 - Consistent performance across all selectivities
- Cassandra:
 - Significantly worse performance with partition key on unique1/unique2
 - Performance more in line with BigQuery with partition key on ten/four/two
 - Likely impacted by how data is spread across cluster, and how many nodes must be visited to perform query
- Actual Results (in contrast with Expected):
 - At higher selectivities, BigQuery outperformed Cassandra significantly
 - At 10% selectivity, Cassandra with partition key on ten/four/two performed best, but was outperformed by BigQuery for higher selectivities

Overall Conclusions

- BigQuery is dominant and EASY
- Cassandra requires careful implementation
 - Partition Key choice impacts spread of data across cluster
 - Desired queries should utilize the clustering effectively
- BigQuery is cheap, while Cassandra requires persistent disk that can be costly in GCP
- Cassandra's optimal performance likely not achieved

Lessons Learned

- Comparisons between systems requires nuance, and brute-force testing doesn't always tell the entire story
- Cassandra punishes the ill-informed
 - Requires intimate knowledge of how it functions to bring out its power
- BigQuery is a great option that requires little investment for big returns



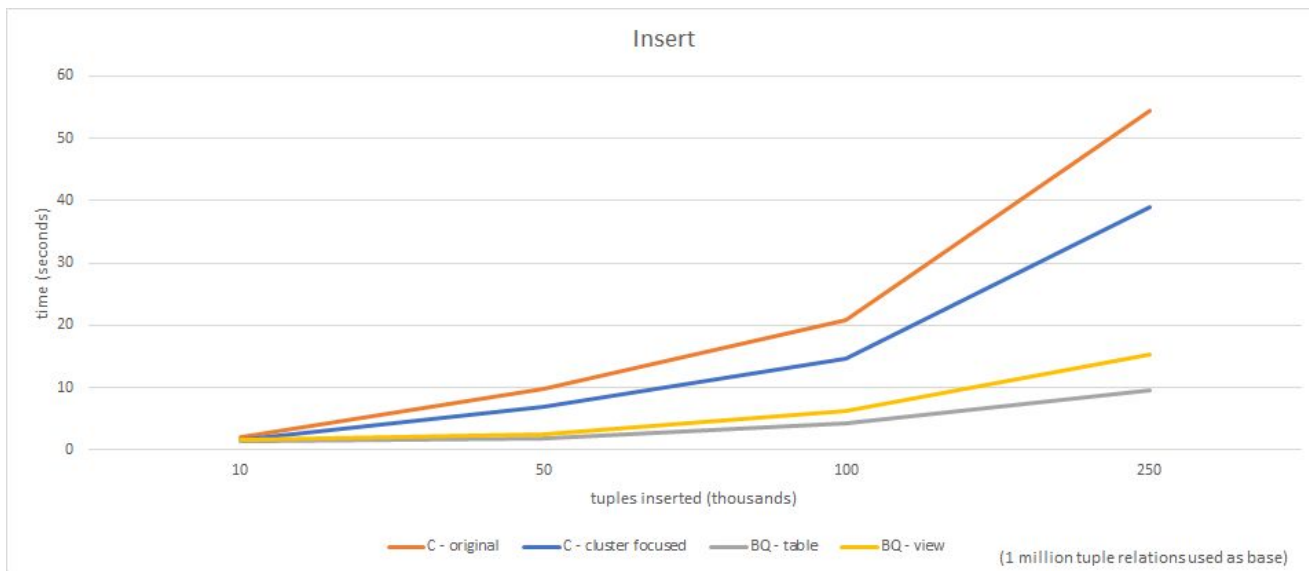
(<https://fictionphile.com/cassandra-of-troy/>)

Appendix

3. Tuple Insertion Experiment

- Designed to test the performance of each database when adding tuples to already existing relations
- Various amounts of tuples (10k, 50k, 100k, 250k) added to relations of 1 million tuples
- General Query Format: `COPY relation_name (attributes)
FROM 'filename.csv' WITH HEADER = true;`
- Expected results:
 - BigQuery will have better performance, as there are no indices to keep updated
 - Cassandra will be slower because there is more internal upkeep with adding tuples (ie: partitioning data, maintaining indices, etc.)

3. Tuple Insertion Experiment



Tuples Inserted	10 k	50 k	100 k	250 k
C - original	2.14	9.77	20.87	54.41
C - cluster focused	1.59	6.98	14.75	39.07
BQ - table	1.47	1.82	4.29	9.47
BQ - view	1.61	2.41	6.25	15.31

- Inserting tuples into BigQuery relation with view was slightly less performant than without view
- Cassandra had similar performance to BigQuery starting out, but deviated more and more as the amount of tuples inserted grew
- Cassandra was faster inserting into the relation that was built with partitioning in mind

3. Tuple Insertion - Conclusions

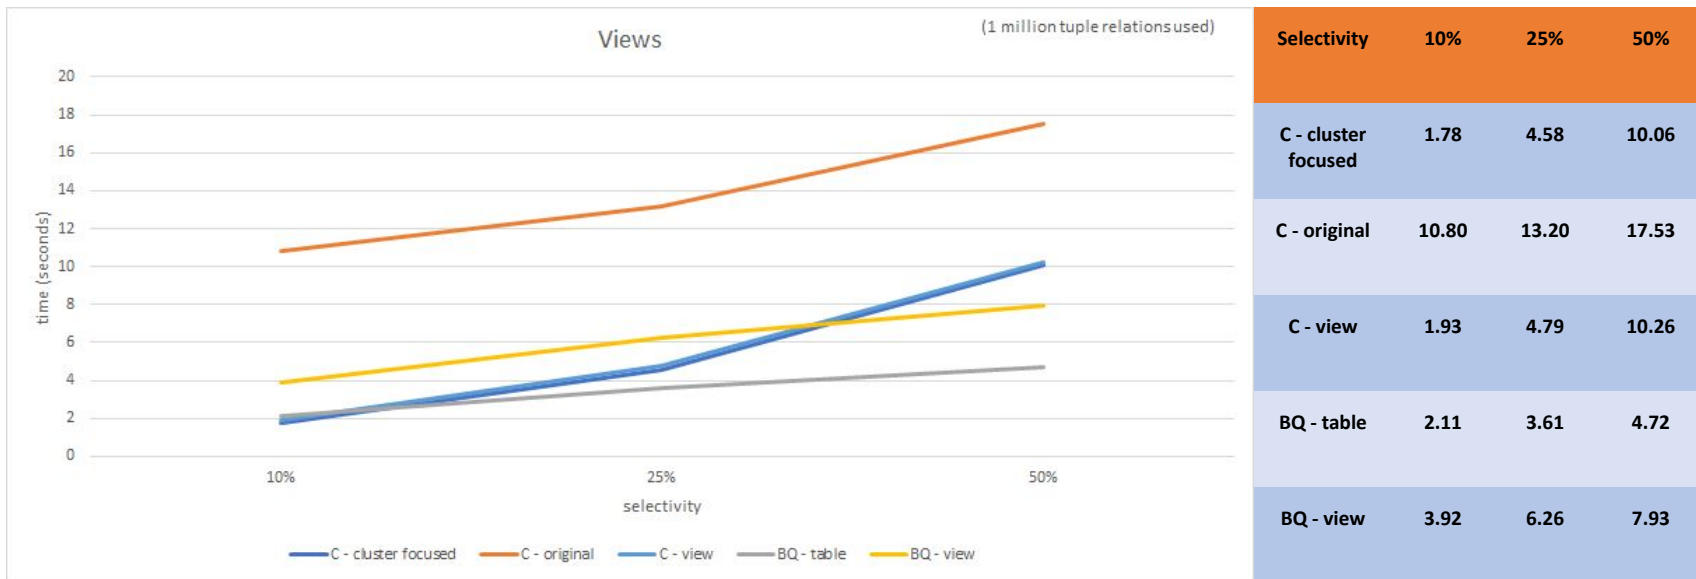
- BigQuery
 - More stable as the number of tuples to insert increased
 - Likely due to not having as much internal maintenance to take care of during insertion (such as maintaining indices, etc).
- Cassandra:
 - Worst performance when inserting tuples into the original relation partitioned on unique2
 - Better performance inserting tuples into the cluster focused relation partitioned on 'ten' attribute
 - Worst case likely had to add all tuples to the same node (since organized by unique2)
 - Better case likely added them more evenly distributed across all nodes in the cluster (since organized by ten), allowing tuples to be added to different nodes concurrently
- Actual Results (in contrast with Expected):
 - BigQuery did have better performance, as expected
 - This is likely attributed to less internal maintenance for BigQuery as opposed to Cassandra (which must partition the data and maintain any indices)

4. Views Experiment

- Designed to test the performance impact using views hason each database
- Used queries with varying selectivities: 10%, 25%, 50%
- General Query Format:

```
SELECT COUNT(unique2)
FROM relation_to_test
WHERE attribute = 0;
(attribute = ten, four, two depending on selectivity)
```
- Expected results:
 - Views will improve performance across the board, but will have greater impact as the relation size grows

4. Views



- BigQuery view hinders performance slightly compared to no view
- Cassandra materialized view improves performance significantly from original relation, and is very comparable to the updated clustered focused relation performance

4. Views Experiment - Conclusions

- BigQuery:
 - Worse performance when querying the view than the relation itself
 - Likely due to overhead required by the view itself
- Cassandra:
 - Much better performance achieved when querying the materialized view than achieved with the original relation partitioned on 'unique2'
 - Materialized view lets you create a version of the relation partitioned on another attribute, which explains why performance achieved matches the performance of the cluster focused relation
- Actual Results (in contrast with Expected):
 - Views did not improve performance of BigQuery, but did improve performance of Cassandra
 - Views didn't make a noticeably greater difference in queries with higher selectivity

Changes to Experiments From Original Design

- Experiment 3: Inserting Tuples into Relations
 - Originally was focused on batch updates
 - Switched, due to limitations encountered in BigQuery, to focus on individual queries that inserted differing numbers of tuples into relations of 1 million tuples
 - Allowed us to attempt to compare impacts of behind the scenes maintenance on performance (such as maintaining indices)
- Experiment 4: Views
 - Originally was designed running the same query on relations of increasing sizes
 - Switched to
 - utilize a consistent 1 million tuple relation
 - varied the selectivity
 - We felt this change would be more interesting to see how views impact performance as the number of tuples returned increases